

Improving Performance

Week 4

COMP6252 (Deep Learning Technologies)
Hikmat Farhat

ECS, University of Southampton

February 19, 2024

- As we have already seen DNN are powerful models.
 - This is mainly due to their strong expressiveness:
 - the ability to model complicated relationships between input and output
 - But this same expressiveness can lead to some problems
- 1 Overfitting
 - Due to their power, NN can "learn" noise present in the training but not in the test datasets
 - This reduces their generalization efficacy
 - 2 Convergence
 - Training takes long time. Sometimes doesn't even converge
 - The results highly dependent on the meta parameters


What can be done?

This week we will introduce three approaches to mitigate the problems of convergence and overfitting.

- 1 Early stopping
- 2 Dropout
- 3 Batch normalization
- 4 Data augmentation

Overfitting

- An example of overfitting is shown below
- The training accuracy keeps increasing while the validation plateau and even decreases




Early stopping

- In the previous figure, the max accuracy occurs at epoch 14
- If we stop training at epoch 14 or close to it not only we get better accuracy but we save lots of time
- This is exactly what is called **early stopping**
- How do we know when to stop training?
- One of the simplest methods is to monitor the **validation loss**


Validation loss vs accuracy

- Below is a plot of validation accuracy vs loss.
- Loss was multiplied by 0.5 to show them both on the same plot
- It is clear that accuracy plateaus after the loss starts increasing




A zoomed view

- Below is a zoomed view of the previous figure
- A vertical line passing by epoch 14 (where max accuracy occurs) is shown




Dropout

- Dropout is another method used to minimize overfitting
Srivastava et al. 2014.
- A dropout layer "drops" certain nodes from the NN as shown
in the figure below



(a) Standard Neural Net



(b) After applying dropout.

From: Srivastava et al.

What is happening exactly?

- Every training pass **different** nodes are dropped
- A neural network with n nodes that uses dropout, can be seen as a set of 2^n different "smaller" networks. Why?
- It means that using dropout is equivalent to averaging over an ensemble of Neural networks sharing the same weights
- Since the noise learned by each "smaller" network is random, "averaging" makes them cancel each other.

Implementation

- Using layers with an on/off switch on nodes is not practical
- More importantly, with an NN with large number of nodes it is **impossible**
- Instead, a node is turned off by multiplying its output with zero **with probability p**
- Mathematically, let y^l be the output of layer l in a fully connected network
- Then the output of layer $(l + 1)$ is given by ($*$ is element wise product)

$$\begin{aligned}r^l &= \text{Bernoulli}(p) \\ \tilde{y}^l &= r^l * y^l \\ z^{l+1} &= W^{l+1} \tilde{y}^l + b^{l+1} \\ y^{l+1} &= \sigma(z^{l+1})\end{aligned}$$

What about testing?

- We have seen that training a NN with dropout is equivalent to training multiple "smaller" networks.
- The result is an average over all such networks.
- To be consistent we must use the weights learned in the training phase for testing
- This means we have to **average** over all possible networks, which is infeasible.
- We approximate this averaging by **not using dropout** during testing
- But the weights obtained in the test/validate phase are multiplied by the probability $1 - p$
- Equivalently the weights are multiplied by $\frac{1}{1-p}$ during training

- A Batch normalization layer transforms the distribution of the input to a distribution that has 0 mean and unit variance
- Batch normalization is a powerful method that
 - 1 Allows for better convergence of NN training
 - 2 Adding BN layers to existing NN yield better generalization results
 - 3 Reduces over fitting
 - 4 Training is less sensitive to meta-parameters (learning rate, weights initialization)

Batch Normalization- How does it work

- Given a mini-batch of tensors x_{ci} of dimension (S,C,H,W)
- where c is the channel index and i collectively refers to all other dimensions.
- Let $N = S \times H \times W$.
- Batch normalization computes the mean and variance of the batch (per channel) according to

$$\mu_c = \frac{1}{N} \sum_{i=1}^N x_{ci}$$
$$\sigma_c^2 = \frac{1}{N} \sum_{i=1}^N (x_{ci} - \mu_c)^2$$

Batch Normalization- How does it work

Given input x_{ci} the normalized inputs of the layer are computed as follows:

$$\hat{x}_{ci} = \frac{x_{ic} - \mu_c}{\sqrt{\sigma_c + \epsilon}}$$

Where ϵ is a small number added for numerical stability. Therefore, for each channel, the $\hat{x}_{ci}$ have zero mean and unit variance.

The output of the batch normalization layer is given by

$$y_{ic} = \gamma_c \hat{x}_{ic} + \beta_c$$

Where γ_c and β_c are **learnable** parameters.

Batch normalisation: inference

- At every step during training, μ_c and σ_c are computed using the mini-batch as described above
- Also, running averages of the mean and variance $\hat{\mu}_c$, $\hat{\sigma}_c$ are computed
- Specifically, at each step when μ_c and σ_c are computed, and the running averages are updated as follows:

$$\hat{\mu}_c \leftarrow 0.9\hat{\mu}_c + 0.1\mu_c$$

$$\hat{\sigma}_c \leftarrow 0.9\hat{\sigma}_c + 0.1\sigma_c$$

- During inference, for input x_{ci} , the output of the batch normalization layer is given by

$$y_{ic} = \frac{x_{ic} - \hat{\mu}_c}{\sqrt{\hat{\sigma}_c}}$$

Batch normalization - Why does it work?

- Considerable ongoing debate
- In the original paper of Ioffe and Szegedy 2015 the claim is that it fixes the internal covariate shift
- meaning the change in the distribution due to the change of the weights during learning
- A different interpretation was given by Kohler et al. 2019: it separates the change in magnitude from the change of directions.
- This can be seen from the fact that the $\hat{x}_{ci}$ are normalized and all of them share the same magnitude γ
- Another recent interpretation can be found in Chai et al. 2020

Example

Consider a mini-batch of size two, containing the tensors of size 2×2 , A and B with both having a single channel.

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

$$B = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$$

- The mean is 4.5 and (biased) variance is 5.25 therefore the output is


$$A = \begin{bmatrix} (1-4.5)/5.25 & (2-4.5)/5.25 \\ (3-4.5)/5.25 & (4-4.5)/5.25 \end{bmatrix}$$

$$B = \begin{bmatrix} (5-4.5)/5.25 & (6-4.5)/5.25 \\ (7-4.5)/5.25 & (8-4.5)/5.25 \end{bmatrix}$$

- When the input has multiple channels, the same operation is performed for each channel independently


- Data: CIFAR10
- The NN has two convolutional blocks
- The first contains two sub-blocks each with
 - 1 32 filters with receptive field of 3×3
 - 2 followed by BN layer
 - 3 followed by ReLU
- the sub-blocks are followed by a max pooling layer of size 2×2
- The second block is the same as the first but with 64 filter
- The conv blocks are followed by two feedforward layers with 128 and 10 nodes respectively

Network Architecture




-  convolutional + BN+ ReLU
-  max pooling
-  fully connected + ReLU
-  fully connected + ReLU

Results: validation accuracy of BN vs no BN




Results: train-vs-validation accuracy with BN



- Batch normalization give much better accuracy.
- Peak value for the accuracy for BN is 78% vs 72%
- The average after the peak for BN is vs 75% vs 69%
- Also, with BN the peak acc is reached after 8 epochs vs 14 without BN.

Results: train-vs-validation accuracy without BN



Data augmentation

- Data augmentation is one method of obtaining higher test/validation accuracy
- Data augmentation is typically done "on-the-fly" in PyTorch
- This is done by applying some transforms to the input
- Example transforms
 - Rotation
 - Translation
 - Horizontal flip
 - Random crop
 -

Example: RandomHorizontalFlip

- Horizontally flip the image with probability 0.5
- RandomHorizontalFlip($p=0.5$)



Example: Translation

- `RandomAffine(degrees=0,translate=(0.3,0.3))`
- Used here for translation only
- Can also perform *rotation*, *scale*, *shear*, etc...



Example: Rotation

- RandomRotation(45)
- Rotate between -45 and 45 degrees



Example: Random crop

1

`RandomCrop((150,150))`



Auto augmentation methods

- AutoAg methods use a policy to transform the input
- For example which transform to use (e.g. rotation, crop) with what probability
- What values are used for the parameters (e.g. angle)
- Some of use search algorithms to determine the policy (e.g. AutoAugment, Cubuk et al. 2018)
- Some are very simple (e.g. TrivialAugment, Müller and Hutter 2021)

SGD with momentum

- A simple variant of SGD to reduce fast changing gradient
- It usually leads to better training convergence
- Let w be the learnable weights and α the learning rate
- Instead of

$$w = w - \alpha \nabla_w \mathcal{L}$$


- use

$$dw = \beta dw + (1 - \beta) \nabla_w \mathcal{L}$$

$$w = w - \alpha dw$$

- where β is a parameter, usually 0.9

Why is it called moving average?



Why is it called exponential moving average?

$$\begin{aligned}dw_t &= \beta dw_{t-1} + (1 - \beta) \nabla_t \mathcal{L} \\&= \beta(\beta dw_{t-2} + (1 - \beta) \nabla_{t-1} \mathcal{L}) + (1 - \beta) \nabla_t \mathcal{L} \\&= \beta^3 dw_{t-3} + \beta^2(1 - \beta) \nabla_{t-2} \mathcal{L} + \beta(1 - \beta) \nabla_{t-1} \mathcal{L} + (1 - \beta) \nabla_t \mathcal{L} \\&= \beta^{t-1}(1 - \beta) \nabla_1 \mathcal{L} + \beta^{t-2}(1 - \beta) \nabla_2 \mathcal{L} + \dots + (1 - \beta) \nabla_t \mathcal{L}\end{aligned}$$


- Usually $\beta < 1$ so $a > b \Rightarrow \beta^a < \beta^b$
- Moreover, $n \rightarrow \infty \Rightarrow \beta^n \rightarrow 0$
- and $\beta = 0 \Rightarrow dw_t = \nabla_t \mathcal{L}$

SGD with momentum

- Validation accuracy vs epoch for ResNet18 with random initial weights

1 SGD with momentum 0.9: pink

2 Vanilla SGD: purple



References

-  Chai, Elaina et al. (2020). “Separating the Effects of Batch Normalization on CNN Training Speed and Stability Using Classical Adaptive Filter Theory”. In: *2020 54th Asilomar Conference on Signals, Systems, and Computers*, pp. 1214–1221.
-  Cubuk, Ekin Dogus et al. (2018). “AutoAugment: Learning Augmentation Policies from Data”. In: *CoRR abs/1805.09501*.
-  Ioffe, Sergey and Christian Szegedy (2015). “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”. In: *CoRR abs/1502.03167*. arXiv: 1502.03167.
-  Kohler, Jonas et al. (2019). “Exponential convergence rates for batch normalization: The power of length-direction decoupling in non-convex optimization”. In: *The 22nd International Conference on Artificial Intelligence and Statistics*. PMLR, pp. 806–815.
-  Müller, Samuel G. and Frank Hutter (2021). “TrivialAugment: Tuning-free Yet State-of-the-Art Data Augmentation”. In: *CoRR abs/2103.10158*.
-  Srivastava, Nitish et al. (2014). “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”. In: *Journal of Machine Learning Research* 15.56, pp. 1929–1958.